

The Authority of Neural Scale

KAFKAS M. CAPRAZLI*, Independent Researcher, Germany

The information retrieval industry is investing billions of dollars to address scaling challenges that may have been solved architecturally in 1985. While modern vector databases struggle to scale beyond 100 million documents, hitting a wall of GPU inefficiency and exorbitant costs, we identify a recurring pattern: successful modern systems—from Meta’s FAISS to Microsoft’s SPANN—have independently converged on architectural principles resembling those of CDS/ISIS, a UNESCO system originally designed for developing nations.

This convergence, we argue, is not coincidental; it may represent an architecturally inevitable response to fundamental limits of information organization. We propose the **Two-Level Theory**, suggesting that efficient retrieval requires separating **semantic understanding** (Level 1) from **structural organization** (Level 2). Modern systems encounter difficulties when they force neural tools to perform symbolic work.

We present the **Neural-to-Symbolic Bridge**, a proposed framework that uses modern AI to populate classical B-tree indices. Our theoretical analysis suggests that replacing brute-force $O(n \cdot d)$ computation with $O(\log k)$ lookups could substantially reduce infrastructure costs and stabilize latency. The future of AI scaling may not require new invention—it may require remembering the optimal architecture we already have.

Additional Key Words and Phrases: information retrieval, vector databases, CDS/ISIS, neuro-symbolic AI, architectural convergence

1 Introduction: The Rediscovery

In 1985, an engineer named Giampaolo Del Bigio faced a seemingly impossible constraint: he had to build a database that could search millions of records on a machine with 64KB of RAM [4]. To survive, he invented an architecture of radical efficiency.

Forty years later, Silicon Valley faces a different constraint: the “Memory Wall” of H100 GPUs. Despite billions in investment, the industry has hit a ceiling. Vector databases are failing at scale [1]. RAG systems are burning cash on idle compute. The technology is different, but the physics are the same.

As a result, the industry appears to be unconsciously migrating back to Del Bigio’s solution.

This article is not a critique of modern AI. It is a map of what we argue is an **inevitable convergence**. We demonstrate that seven of the world’s most advanced systems—built by Microsoft, Google, and Meta—have unknowingly recreated specific architectural mechanisms resembling those of UNESCO’s 1985 software, CDS/ISIS.

This convergence, we suggest, is not mere technological coincidence. It may be a mathematical inevitability arising from fundamental limits of information organization [18]. Hidden in UNESCO archives lies documentation for a system that achieved what billion-dollar companies now struggle to replicate: $O(\log n)$ retrieval complexity through hierarchical B-tree indexing and variable-length encoding [15].

Modern systems aren’t innovating in isolation. They appear to be converging—unconsciously, expensively, perhaps inevitably—on an architecture that poverty necessitated and mathematics validated forty years ago. This paper presents the evidence for this convergence, the theory explaining why it may be necessary, and a framework that bridges the two worlds.

It is time we stopped inventing, and started understanding.

*This manuscript represents a novel form of human-AI collaborative research. The author extensively collaborated with multiple AI systems (Claude Opus 4.1, Perplexity, Gemini Pro) for literature synthesis and drafting. The core insight—that modern systems are unconsciously converging on CDS/ISIS principles—represents the author’s original contribution. The author takes full responsibility for all claims.

2 The Two-Level Theory of Information Retrieval

The crisis in scalable information retrieval stems from what we characterize as a category error: attempting to solve a two-level problem with a single-level solution. We propose that retrieval naturally decomposes into two distinct layers, and modern systems struggle when they ignore forty years of evidence that efficient, billion-scale access requires architectural decomposition.

2.1 Level 1: Semantic Understanding (The Neural Domain)

Modern neural embeddings represent humanity’s greatest achievement in computational semantics. Transformer models map arbitrary text into high-dimensional vector spaces $\mathbb{R}^d$ where geometric distance correlates with semantic similarity. A legal contract about wheat futures and an agricultural report on grain diseases—meaningless to keyword matching—cluster naturally in embedding space.

This brilliance has an unavoidable limit. Information retrieval requires finding k -nearest neighbors (k -NN) among n documents, each a d -dimensional vector. The brute-force baseline requires calculating similarity $S(q, d_i)$ between query q and every document d_i . With n documents and d dimensions, complexity is $O(n \cdot d)$ —linear in both collection size and embedding dimension.

The mathematics are unforgiving. With typical $d = 768$ and $n = 1$ billion, each query requires 768 billion floating-point operations. No GPU parallelization changes this fundamental complexity; it merely distributes the computation. Modern systems implicitly acknowledge this: efficiency gains consistently come from structural processing—Level 2.

2.2 Level 2: Structural Organization (The Symbolic Domain)

Level 2 addresses access—transforming search from comparison to lookup. Forty years ago, facing 64KB memory constraints, CDS/ISIS engineers discovered what Silicon Valley is rediscovering: efficient retrieval requires symbolic organization, not semantic computation.

UNESCO’s CDS/ISIS solved Level 2 through two interconnected structures:

- (1) **Inverted File (IF) Indexing:** Instead of searching data, search an index of discrete terms linked to documents via posting lists [20]. This transforms content search into pointer traversal.
- (2) **B-tree Hierarchical Structure:** The index itself is organized as a balanced tree, guaranteeing that any lookup requires at most $\log_m(k)$ node traversals, where k is the index size and m is the branching factor.

The mathematical elegance is striking: for a billion documents indexed into $k = 10,000$ clusters with $m = 100$, we need only $\log_{100}(10,000) = 2$ operations versus one billion comparisons.

2.3 The Fundamental Incompatibility

The incompatibility between approaches appears mathematical. Consider retrieval complexity:

$$C_{total} = C_{semantic} + C_{structural}$$

Pure neural systems minimize $C_{semantic}$ through better embeddings while setting $C_{structural} \approx 0$, forcing all complexity into the semantic layer: $C_{total} \approx O(n \cdot d)$. No optimization overcomes this linear scaling. Pure symbolic systems (pre-2017) minimize $C_{structural}$ through hierarchical indexing ($O(\log k)$) but fail on semantic variation.

Information theory suggests a decomposition necessity [3]. For scalability, the system must likely achieve:

- $C_{semantic} = O(d)$ for local embedding computation

- $C_{structural} = O(\log k)$ for hierarchical retrieval
- $C_{total} = O(d + \log k) \approx O(\log k)$ since d is fixed.

3 The Convergence Evidence: Seven Systems, One Pattern

The most compelling evidence for our theory comes from an unexpected source: the R&D departments of the world’s leading technology companies. Seven independent teams have unknowingly recreated solutions resembling those that UNESCO distributed free to developing nations in 1985.

The pattern is unmistakable. Each system, developed in isolation, converges on similar architectural principles:

Table 1. The Convergence: Modern Systems Rediscovering CDS/ISIS Principles

System	Year	“Innovation”	Analogous Principle	Mapping
FAISS IVF [8]	2017	Inverted files for vectors	Inverted File (IF) Indexing	Posting lists for centroids
DiskANN [14]	2019	Hierarchical navigable graph	Hierarchical Navigation	Logarithmic graph traversal
ScaNN [6]	2020	Anisotropic quantization	Variable-Length Encoding	Adaptive bit allocation
SPANN [2]	2021	Hierarchical balanced clustering	Hierarchical Decomposition	Orthogonal indexing
Continuous Batching [12]	2024	Ragged batching, KV cache	Variable Records	No-padding storage
Pinecone [11]	2024	Serverless vector index	Distributed IF	Web-scale inverted files
Weaviate [16]	2025	Hybrid HNSW+inverted	Dual Index Architecture	Graph + IF combination

Consider Meta’s FAISS. In 2017, it introduced “Inverted Files for vectors” (IVF), partitioning vectors into Voronoi cells. This approach closely resembles CDS/ISIS’s inverted file structure from 1985. The key modification: where CDS/ISIS indexed text tokens, FAISS indexes vector centroids.

Microsoft’s DiskANN achieves billion-scale search through a “Vamana graph”—a hierarchical structure. The paper’s abstract announces: “hierarchical navigation enables logarithmic search complexity.” This insight appears almost verbatim in Chapter 3 of the 1985 CDS/ISIS manual [15].

Unsatisfied with DiskANN, Microsoft developed SPANN in 2021, introducing “hierarchical balanced clustering”—an approach that mirrors what CDS/ISIS documentation describes as field-independent hierarchical organization. The reported results are striking: 96% recall on billion-scale datasets with significant memory cost reductions [2].

These may not be independent discoveries—they may be forced convergences. Information retrieval faces immutable constraints that create what we term a “solution funnel.” Semantic understanding requires continuous representations. Structural organization requires discrete hierarchies. Economic reality prohibits $O(n)$ scaling.

4 Historical Vindication: Del Bigio and Rybiński

To understand why the future of information retrieval may look exactly like 1985, we must look at the constraints that forged the past.

4.1 Del Bigio’s Constraint-Driven Genius

The primary architect of CDS/ISIS, Giampaolo Del Bigio, was not optimizing for data centers. He was designing for libraries in Lagos and Manila using PDP-11s and early PCs with 64KB RAM. In this environment, inefficiency was a hard stop.

Del Bigio responded with radical efficiency. He implemented variable-length field encoding (ISO 2709), what some consider a direct ancestor of the ragged tensors now used in continuous batching [7]. He architected the system’s “Master File” (storage) and “Inverted File” (access) as independent entities, bridged only by minimal pointers. His design demonstrated that hierarchical decomposition was practically mandatory to minimize disk seek times—a constraint that mirrors the memory bandwidth bottlenecks facing H100 GPUs today.

4.2 Rybiński and the Web-Scale Proof

If Del Bigio proved the architecture on small machines, Henryk Rybiński proved it at scale. In the 1990s, Rybiński adapted the CDS/ISIS engine for the internet (WWW-ISIS), powering the FAO’s agricultural database (AGRIS) [13]. Rybiński demonstrated that the Inverted File + B-tree architecture was natively parallelizable, achieving high-concurrency operations long before “serverless” became a buzzword.

4.3 The Innovation Paradox

Why did Silicon Valley overlook these solutions? The answer may lie in what we call the **Innovation Paradox**: abundance breeds inefficiency, while scarcity breeds optimality. For decades, Moore’s Law allowed engineers to rely on brute force. But as data volume exploded to the trillion-token scale, we entered a new era of constraint—not of disk space, but of GPU memory and energy. We may not be witnessing a new invention; we may be witnessing the industry migrating from the Architecture of Abundance back to Del Bigio’s Architecture of Necessity.

5 The Solution: Neural-to-Symbolic Bridge

We have established that scalable retrieval likely requires satisfying two opposing constraints. We propose a unified conceptual framework: the **Neural-to-Symbolic Bridge**. This architecture proposes utilizing modern neural networks solely to generate the “keys” for a classical, mathematically efficient indexing “engine.”

5.1 The Framework Architecture

The architecture operates as a feed-forward pipeline:

$$\text{Query} \xrightarrow{\text{LLM}} \text{Embedding } (\mathbb{R}^d) \xrightarrow{\text{Quantization}} \text{Symbol } (\Sigma) \xrightarrow{\text{B-Tree}} \text{Retrieval}$$

- (1) **Neural Projection (Level 1)**: An LLM maps text into a dense vector $v \in \mathbb{R}^d$.
- (2) **Hierarchical Quantization (The Bridge)**: We apply a function $Q(v) \rightarrow s$, mapping the continuous vector to a discrete symbol s (centroid ID) [19]. Functionally, this reintroduces the principle of **Unified Authority Files** [17]: constraining infinite semantic variance into finite, canonical identifiers to promote retrieval stability.
- (3) **Symbolic Indexing (Level 2)**: The discrete symbol s serves as the search key for a standard, inverted-file B-tree index.

5.2 Projected Advantages

By shifting the burden to Level 2, the framework is predicted to yield structural gains:

- **Complexity Reduction:** The architecture suggests a reduction in complexity from $O(n \cdot d)$ to $O(\log k)$.
- **Cost Reduction:** This eliminates the need for full-precision vectors in GPU memory. Architectural analyses of efficient RAG implementations suggest that moving inverted lists to SSDs could substantially reduce infrastructure costs for comparable recall [10].
- **Latency:** By removing the GPU bottleneck from the retrieval path, latency is expected to stabilize in the single-digit millisecond range.

5.3 Theoretical Validation & Industry Parallels

The viability of this architecture is supported by the very systems that have unintentionally adopted its principles.

- **Theoretical Validation:** Our model aligns with the **Information Bound Theorem**, suggesting that decomposing semantic entropy from structural entropy may be the only path to maintain recall at $O(\log k)$ [9].
- **Derived Empirical Evidence:** Microsoft’s SPANN implements this exact Level 2 “field independence” principle, reporting 96% recall on billion-scale datasets with substantial memory cost reductions [2]. While we have not independently verified these commercial gains, the reported improvements align with our theoretical predictions.
- **Historical Proof:** The scalability of this stack (Inverted Files + B-trees) was validated at scale by AGRIS in the 1990s [5].

6 Implications and Future: The Recursive Horizon

The convergence on CDS/ISIS principles may serve as a roadmap for the future.

Democratizing Semantic Search: Today’s cost structure represents an artificial barrier. Our framework could correct this. A law firm in Mumbai or a research station in Kenya could index their history on commodity hardware. By shifting from $O(n \cdot d)$ compute to $O(\log k)$ lookups, we may democratize access to semantic search, potentially redirecting computational resources toward actual research.

The Recursive Frontier: While two levels solve today’s challenges, the future may be recursive.

- Level 2: Documents $\rightarrow$ Symbols ($O(\log k)$)
- Level 3: Symbols $\rightarrow$ Meta-patterns ($O(\log \log k)$)
- Level 4: Patterns $\rightarrow$ Knowledge structures ($O(\log \log \log k)$)

At trillion-scale, recursive decomposition may become mandatory. This anticipates the future just as Del Bigio anticipated ours.

7 Conclusion: The Inevitable Architecture

The scaling crisis in information retrieval may not be a failure of technology—it may be a failure of memory. For a decade, the industry has attempted to solve a structural problem with semantic tools, fighting mathematical certainty.

Our analysis of seven modern systems reveals what appears to be a unified truth: all scalable systems eventually converge on the principles of **hierarchical decomposition**, **field independence**, and **variable-length encoding**. These are not modern inventions. They are the core tenets of UNESCO’s CDS/ISIS.

This convergence suggests that these principles may be **fundamental constants** of information organization. The path forward is the **Neural-to-Symbolic Bridge**: using the semantic brilliance of modern AI to populate the structural efficiency of classical indexing. Rigorous mathematical validation and implementation details will be presented in a forthcoming technical paper, but the convergence pattern alone demands immediate attention.

The architecture for the future may have been waiting in UNESCO archives for forty years.

CDS/ISIS: The Forgotten Giant

- **Scale**: Designed to handle millions of records on 64KB RAM (1985).
- **Reach**: Deployed in over 140 countries, serving millions of users.
- **Innovations**: Variable-length fields, inverted files, hierarchical indexing.
- **Lesson**: Constraint-driven design often achieves mathematical efficiency.

Key Takeaways

- Modern IR systems (FAISS, SPANN, DiskANN) appear to be unconsciously recreating 1985 CDS/ISIS architecture.
- The proposed **Two-Level Theory** explains why: Semantic understanding ($O(n \cdot d)$) must be separated from structural organization ($O(\log k)$).
- Our **Neural-to-Symbolic Bridge** framework synthesizes these principles to project substantial cost reduction.
- Constraint is often the mother of optimal architecture; we may need to look to “legacy” systems for future scaling.

References

- [1] Tilmann Bruckhaus. 2024. *RAG does not work for enterprises*. Technical Report. Strative Industry Report. <https://arxiv.org/pdf/2406.04369.pdf>
- [2] Qi Chen, Bing Wang, Mingqin Zhang, et al. 2021. SPANN: Highly-efficient billion-scale approximate nearest neighbor search. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 34. https://proceedings.neurips.cc/paper_files/paper/2021/hash/299dc35e747eb77177d9cea10a802da2-Abstract.html
- [3] T. Y. Chen et al. 2025. Rate-Distortion Limits for Multimodal Retrieval: Theory, Optimal Codes, and Finite-Sample Guarantees. *arXiv preprint arXiv:2509.11054* (2025). <https://arxiv.org/abs/2509.11054> Columbia University.
- [4] Giampaolo Del Bigio. 1995. *The CDS/ISIS System: Technical Foundations*. Technical Report. UNESCO, Paris. <https://unesdoc.unesco.org/ark:/48223/pf0000101570>
- [5] FAO, Rome (Italy). Library and Documentation Systems Div. 2005. *The AGRIS Application Profile for the International Information System on Agricultural Sciences and Technology: Guidelines on Best Practices for Information Object Description*. Technical Report. Food and Agriculture Organization of the United Nations, Rome, Italy. <https://openknowledge.fao.org/server/api/core/bitstreams/00b4c742-6673-4420-bb66-5b7c2ba3a668/content/ae909e01.htm> Version 2.0, 81 p..
- [6] Ruiqi Guo, Philip Sun, Erik Lindgren, et al. 2020. Accelerating Large-Scale Inference with Anisotropic Vector Quantization. In *International Conference on Machine Learning (ICML)*. 3887–3896. <http://proceedings.mlr.press/v119/guo20c.html>
- [7] International Organization for Standardization. 2008. ISO 2709:2008 Information and documentation—Format for information exchange. International Standard. <https://www.iso.org/standard/41319.html>
- [8] Jeff Johnson, Matthijs Douze, and Hervé Jégou. 2017. Billion-scale similarity search with GPUs. *arXiv preprint arXiv:1702.08734* (2017). <https://arxiv.org/abs/1702.08734>
- [9] Omri Korten, T. Pitassi, and R. Impagliazzo. 2025. Stronger Cell Probe Lower Bounds via Local PRGs. *Electronic Colloquium on Computational Complexity (ECCC)* TR25-030 (2025). <https://eccc.weizmann.ac.il/report/2025/030/>

- [10] Market Research Future. 2024. *Vector Database Market Report*. Technical Report. Industry Analysis Quarterly. <https://www.marketresearchfuture.com/reports/vector-database-market-32105> Accessed: 2025-12-01.
- [11] Pinecone Systems. 2024. Pricing and Performance Documentation. Technical Documentation. <https://www.pinecone.io/pricing/> Accessed: 2025-12-05.
- [12] Rémi Ouazan Reboul, Arthur Zucker, and Luc Georges. 2025. Continuous Batching: From Attention to High-Throughput Inference. Hugging Face Technical Blog. https://huggingface.co/blog/continuous_batching Accessed: 2025-12-09.
- [13] Henryk Rybiński. 1999. WWW-ISIS: Scaling CDS/ISIS to the Web. *ICIE Technical Report* (1999). <http://w2isis.icml9.org/activity.php?lang=en&id=33>
- [14] Suhas Subramanya, F. Devvrit, Harsha Vardhan Simhadri, Ravishankar Krishnawamy, and R. Kadekodi. 2019. DiskANN: Fast accurate billion-point nearest neighbor search on a single node. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 32. <https://proceedings.neurips.cc/paper/2019/hash/093f65e080a295f8076b1c5722a46aa2-Abstract.html>
- [15] UNESCO. 1989. *CDS/ISIS Reference Manual Version 3.0*. UNESCO, Paris. <https://unesdoc.unesco.org/ark:/48223/pf0000194203>
- [16] Weaviate. 2025. Hybrid Search Architecture. Technical Blog. <https://weaviate.io/blog/hybrid-search-explained> Accessed: 2025-12-05.
- [17] James Weinheimer and Kafkas Caprazli. 2003. A Framework for Unified Authority Files: A Case Study of Corporate Body Names in the FAO Catalogue. In *Research and Advanced Technology for Digital Libraries*, Traugott Koch and Ingeborg Torvik Sølberg (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 374–386. doi:10.1007/978-3-540-45175-4_34
- [18] Orion Weller, Michael Boratko, I. Naim, and J. Lee. 2025. On the theoretical limitations of embedding-based retrieval. *arXiv preprint arXiv:2508.21038* (2025). <https://arxiv.org/abs/2508.21038>
- [19] Ali Zandieh, M. Daliri, M. Hadian, and V. Mirrokni. 2025. *TurboQuant: Online Vector Quantization with Near-optimal Distortion Rate*. Technical Report. Google Research. <https://arxiv.org/abs/2504.19874>
- [20] Justin Zobel and Alistair Moffat. 2006. Inverted files for text search engines. *Comput. Surveys* 38, 2 (2006), 6. doi:10.1145/1132956.1132959

Acknowledgments

Kafkas M. Caprazli is an independent researcher who was part of FAO teams that implemented CDS/ISIS-based information systems including AGRIS, CARIS, ASFA, and DOCREP. He contributed to multilingual implementations including Thai AGROVOC. He co-authored research on unified authority files [17] and currently studies the convergence between historical and modern information retrieval architectures.